

Intersects) look for Records in common
 join . . . key

Join : ① Inner join (USING)

② Case self join
 ③ Left join . (most common) <> Right join

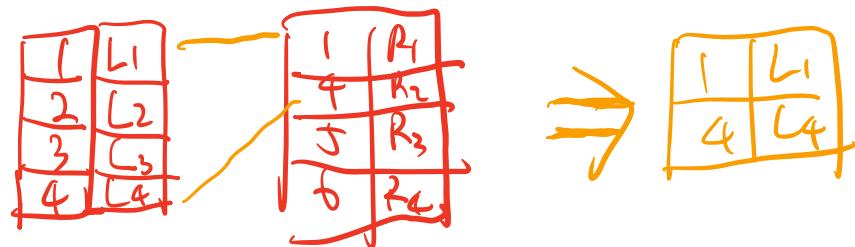
⑤ Full join

⑥ CROSS join

Union (all)

Intersect

⑦ semi join



⑧ anti-join

Exercise

Relating semi-join to a tweaked inner join

Let's revisit the code from the previous exercise, which retrieves languages spoken in the Middle East.

```
SELECT DISTINCT name
FROM languages
WHERE code IN
  (SELECT code
   FROM countries
   WHERE region = 'Middle East')
ORDER BY name;
```

Sometimes problems solved with semi-joins can also be solved using an inner join.

```
SELECT languages.name AS Language
FROM languages
INNER JOIN countries
ON languages.code = countries.code
WHERE region = 'Middle East'
ORDER BY language;
```

This inner join isn't quite right. What is missing from this second code block to get it to match with the correct answer produced by the first block?

Subquery made WHERE clause.

Types of joins

- INNER JOIN
 - Self-joins
- OUTER JOIN
 - LEFT JOIN
 - RIGHT JOIN
 - FULL JOIN
- CROSS JOIN
- Semi-join / Anti-join

INNER JOIN vs LEFT JOIN

left_table		right_table	
id	val	id	val
1	L1	1	R1
2	L2	4	R2
3	L3	5	R3
4	L4	6	R4



INNER JOIN

L.id	L.val	R.val
1	L1	R1
4	L4	R2

left_table		right_table	
id	val	id	val
1	L1	1	R1
2	L2	4	R2
3	L3	5	R3
4	L4	6	R4



LEFT JOIN

L.id	L.val	R.val
1	L1	R1
2	L2	
3	L3	
4	L4	R2

RIGHT JOIN vs FULL JOIN

left_table		right_table	
id	val	id	val
1	L1	1	R1
2	L2	4	R2
3	L3	5	R3
4	L4	6	R4

left_table		right_table	
id	val	id	val
1	L1	1	R1
2	L2	4	R2
3	L3	5	R3
4	L4	6	R4



RIGHT JOIN

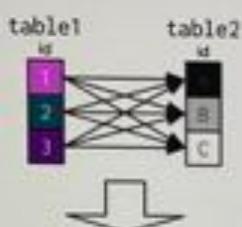
R_id	L_val	R_val
1	L1	R1
4	L4	R2
5		R3
6		R4



FULL JOIN

L_id	R_id	L_val	R_val
1	1	L1	R1
2		L2	
3		L3	
4	4	L4	R2
	5		R3
	6		R4

CROSS JOIN with code

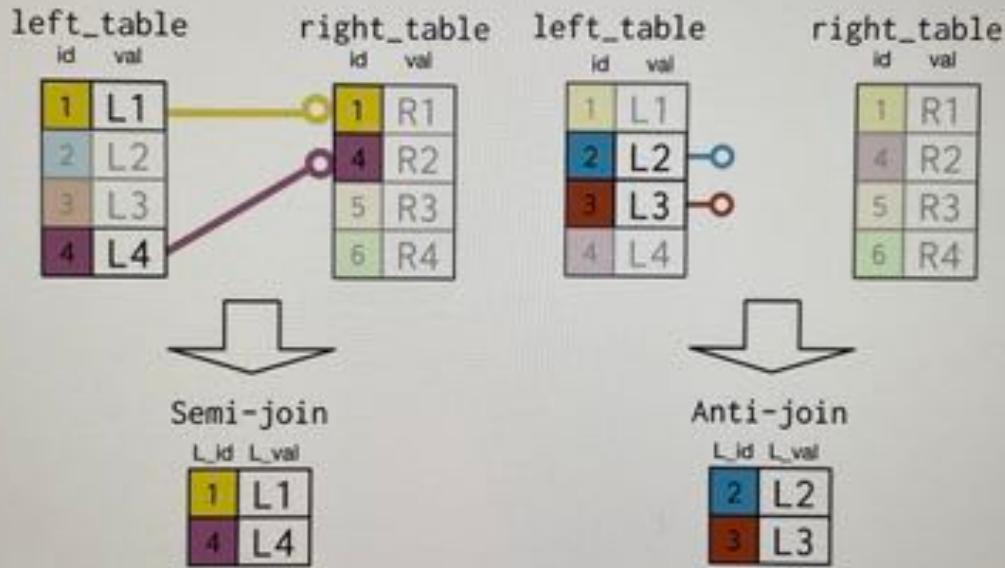


CROSS JOIN

id1	id2
1	
1	B
1	C
2	
2	B
2	C
3	
3	B
3	C

```
SELECT table1.id AS id1,
       table2.id AS id2
  FROM table1
CROSS JOIN table2;
```

Semi-joins and Anti-joins



Final challenge

Welcome to the end of the course! The next three exercises will test your knowledge of the content covered in this course and apply many of the ideas you've seen to difficult problems. Good luck!

Please carefully read the instructions and solve them step-by-step, thinking about how the different clauses work together.

In this exercise, you'll need to get the country names and other 2015 data in the `economics` table and the `countries` table for Central American countries with an official language.

Instructions SQL

- Select unique country names. Also select the total_investment and imports fields.
- Use a left join with `languages` on the left. (An inner join would also work, but please use a left join here.)
- Match on `code`. In the two tables (B), use a subquery inside of (B) to choose the appropriate `languages` records.
- Order by country name ascending.
- Use table aliasing but not field aliasing in this exercise.

Take Hint (10 XP)

```
query.sql
1 SELECT DISTINCT e.name, e.total_investment, e.imports
2 FROM economics AS e
3 INNER JOIN countries AS c
4 USING (code)
5 WHERE year = 2015 AND code IN (
6   SELECT code FROM countries
7   INNER JOIN languages
8   USING (code)
9   WHERE countries.region = 'Central America' AND languages.official = TRUE)
10
```

I

Run Code

query result	economics	languages	countries
name	total_investment		Imports
Belize	22.014		6.743
Costa Rica	20.219		4.829
El Salvador	13.963		3.192
Guatemala	13.453		16.124
Honduras	24.855		9.363

Showing 7 out of 7 rows

两个结果一样

The screenshot shows a SQL exercise interface. On the left, there's a code editor with a partially completed query:

```
SELECT DISTINCT name, total_investment, imports
FROM ...
LEFT JOIN ...
ON L.code = E.code
AND L.code IN (
    SELECT L.code
    FROM Languages AS L
    WHERE official = 'true'
)
-- Where region and year are correct
WHERE region = 'Central America' AND year = 2013
-- Order by field
ORDER BY name;
```

A modal window at the bottom asks "Did you find this hint helpful?" with "Yes" and "No" buttons.

On the right, the interface displays the results of the executed query:

country	economics	languages	countries
Belize	22.04	0.743	
Costa Rica	20.29	4.829	
El Salvador	15.963	6.185	
Guatemala	15.453	15.124	
Honduras	24.653	9.363	

At the bottom, it says "Showing 7 out of 7 rows". A "Run Code" button is visible in the top right of the results area.

Final challenge (2)

What's that was challenging, huh?

Let's clean up a bit and calculate the average fertility rate for each region in 2015.

Instructions

- Include the name of region, its continent, and average fertility rate (elided as avg_fert_rate).
- Sort based on avg_fert_rate ascending.
- Remember that you'll need to GROUP BY all fields that aren't included in the aggregate function of SELECT.

Take Hint (0.00 XP)

Incorrect Submission

Check the first entry in the left operand of the JOIN expression. The checker expected to find countries AS b in there.

Did you find this feedback helpful? Yes No

```
query_id
1 -- Select fields
2 SELECT region, continent, AVG(fertility_rate) AS avg_fert_rate
3   -- From left table
4   -- Join to right table
5   FROM populations AS p
6     -- INNER JOIN countries AS c
7       -- Match on country_code = c.code
8     -- Where specific records matching some condition
9 WHERE year = 2015
10   -- Group appropriately
11 GROUP BY region, continent
12   -- Order appropriately
13 ORDER BY avg_fert_rate;
14
15
16
17
18
```

query result	region	continent	avg_fert_rate
Southern Europe	Europe	1.4261000037165	
Eastern Europe	Europe	1.4904669032275	
Baltic Countries	Europe	1.6033333425985	
Western Europe	Europe	1.8326000037446	
Eastern Asia	Asia	1.83166666256125	

Showing 5 out of 5 rows

Final challenge (3)

Welcome to the last challenge problem. By now you're a query master! Remember that these challenges are designed to take you to the limit to solidify your SQL knowledge! Here's a deep breath and solve this step-by-step.

You are now tasked with determining the top 10 capital cities in Europe and the Americas in terms of a calculated percentage (avg_city.pop / avg.metroarea.pop) in descending order.

Do not use table aliasing in this exercise.

Instructions

- Select the city name, country code, city proper population, and metro area population.
- Calculate the percentage of metro area population composed of city proper population for each city in (USA), ordered by (USA) in descending order.
- Filter only on capital cities in Europe and the Americas in a subquery.
- Make sure to exclude records with missing data on metro-area population.
- Order the result by (USA) in descending order.
- Then determine the top 10 capital cities in Europe and the Americas in terms of this **avg_pct** percentage.

Take Hint (0.00 XP)

Incorrect Submission

We expected to find a column named `country_code` in the result of your query, but you didn't.

Did you find this feedback helpful? Yes No

```
query_id
1 SELECT name, country_code, city_proper_pop, metroarea_pop, city_proper_pop.metroarea_pop*100 AS city_pct
2 FROM cities
3 WHERE metroarea_pop IS NOT NULL AND name IN (
4   SELECT capital FROM countries WHERE continent IN ('Europe' || continent LIKE 'Americas')
5 ORDER BY city_pct DESC
6 LIMIT 10
7 )/SELECT
8   -- Calculate city_pct
9   -- From appropriate table
10  FROM ...
11  -- Where ...
12  WHERE ...
13  WHERE ...
14  -- Subquery
15  (SELECT capital
16    FROM ...
17    WHERE ...
18    OR ...
19    AND ...
20  )/ORDER BY ...
21  -- Order appropriately
22  ORDER BY ...
23  -- Limit result
24
```

query result	name	country_code	city_proper_pop	metroarea_pop	city_pct
Lima	PER	8482000	10760000	80.541651059979	
Bogota	COL	7672783	9800000	80.395744216791	
Moscow	RUS	12100000	16700000	75.45549296052348	
Vienna	AUT	1843861	2800000	71.58875291665932	
Montevideo	URY	1508080	1947654	70.0046154845312	

Showing 10 out of 10 rows

Simple vs. correlated subqueries

Simple Subquery

- Can be run *independently* from the main query
- Evaluated once in the whole query

Correlated Subquery

- *Dependent* on the main query to execute
- Evaluated in loops
 - Significantly slows down query runtime

case when
CTE

Common Table Expressions

Common Table Expressions (CTEs)

- Table *declared* before the main query
- *Named* and *referenced* later in **FROM** statement

Setting up CTEs

```
WITH cte AS (
    SELECT col1, col2
    FROM table)
SELECT          *
    AVG(col1) AS avg_col
FROM cte;
```

Why use CTEs?

- Executed once
 - CTE is then stored in memory
 - Improves query performance
- Improving organization of queries
- Referencing other CTEs
- Referencing itself (`SELF JOIN`)

Differentiating Techniques

Joins

- Combine 2+ tables
 - Simple operations/aggregations

Correlated Subqueries

- Match subqueries & tables
 - Avoid limits of joins
 - High processing time

Multiple/Nested Subqueries

- Multi-step transformations
 - Improve accuracy and reproducibility

Common Table Expressions

- Organize subqueries sequentially
- Can reference other CTEs

Different use cases

Joins

- 2+ tables (*What is the total sales per employee?*)

Correlated Subqueries

- *Who does each employee report to in a company?*

Multiple/Nested Subqueries

- *What is the average deal size closed by each sales representative in the quarter?*

Common Table Expressions

- *How did the marketing, sales, growth, & engineering teams perform on key metrics?*

Exercise

Get team names with a subquery

Let's solve a problem we've encountered a few times in this course so far -- How do you get both the home and away team names into one final query result?

Out of the 4 techniques we just discussed, this can be performed using subqueries, correlated subqueries, and CTEs. Let's practice creating similar result sets using each of these 3 methods over the next 3 exercises, starting with subqueries in #SQL.

Instructions 2/2 Run SQL

- Add a second subquery to the `SELECT` statement to get the away team name, changing only the `homeTeam`. Left join both subqueries to the `match` table on the `id` column.

Warning: if your code is timing out, you have probably made a mistake in the `SELECT` and tried to join on the wrong fields which caused the table to be too big! Read the provided code and comments carefully, and check your `ON` conditions!

Incorrect Submission

Check the second entry in the target list of the `SELECT` statement. The checker expected to find `homeTeam` in there.

Did you find this feedback helpful? Yes No

query.sql

```

1 SELECT
2   m.date,
3     -- Get the home and away team names
4     hometeam,
5     awayteam,
6     m.home_goal,
7     m.away_goal
8   FROM match AS m
9
10  -- Join the home subquery to the match table
11  LEFT JOIN (
12    SELECT match_id, team.team_long_name AS hometeam
13    FROM match
14    LEFT JOIN team
15      ON match.homeTeam_id = team.team_api_id) AS home
16    ON home.id = m.id
17
18  -- Join the away subquery to the match table
19  LEFT JOIN (
20    SELECT match_id, team.team_long_name AS awayteam
21    FROM match
22    LEFT JOIN team
23      -- Get the away team ID in the subquery
24      ON match.awayTeam_id = team.team_api_id) AS away
25    ON away.id = m.id;

```

query result

date	hometeam	awayteam
2011-07-29	Oud-Heverlee Leuven	RSC Anderlecht
2011-07-30	RAEC Mons	Standard de Liège
2011-07-30	KRC Genk	Beerschot AC

Showing 100 out of 1000 rows

datacamp

Exercise

Get team names with correlated subqueries

Let's solve the same problem using correlated subqueries -- how do you get both the home and away team names into one final query result?

This can easily be performed using correlated subqueries, but how might that impact the performance of your query? Complete the following steps and let's find out!

Please note that your query will run more slowly than the previous exercise!

Instructions 2/2 Run SQL

- Create a second correlated subquery in `SELECT`, getting the away team's name.
- Select the home and away goal columns from `match` in the main query.

query.sql

```

1 SELECT
2   m.date,
3     SELECT team.long_name
4     FROM team AS t
5     WHERE t.team_api_id = m.homeTeam_id AS hometeam,
6       -- Connect the team to the match table
7     SELECT team.long_name
8     FROM team AS t
9     WHERE t.team_api_id = m.awayTeam_id AS awayteam,
10      -- Select home and away goals
11      home_goal,
12      away_goal,
13    FROM match AS m;

```

query result

date	hometeam	awayteam	home_goal	away_goal
2011-07-29	Oud-Heverlee Leuven	RSC Anderlecht	2	1
2011-07-30	RAEC Mons	Standard de Liège	1	1
2011-07-30	KRC Genk	Beerschot AC	3	1
2011-07-30	KAAS Gent	KSV Cercle Brugge	0	1
2011-07-30	Sparta Rotterdam	SV Zulte-Waregem	0	0

Showing 100 out of 1000 rows

Exercise

Get team names with CTEs

You've now explored two methods for answering the question. How did you get both the home and away team names into one final query result?

Let's explore the final method - common table expressions. Common table expressions are similar to the subquery method for generating results, mostly differing in syntax and the order in which information is processed.

Instructions 2/2 Run SQL

- Let's declare the second CTE, `away`. Join it to the first CTE on the `id` column.
- The `date`, `home_goal`, and `away_goal` columns have been added to the CTEs. `SELECT` them into the main query.

query.sql

```

1 WITH home AS (
2   SELECT m.id,
3     m.date,
4     t.team_long_name AS hometeam, m.home_goal,
5   FROM match AS m
6   LEFT JOIN team AS t
7     ON m.homeTeam_id = t.team_api_id,
8     -- Declare and set up the away CTE
9   away AS (
10   SELECT m.id,
11     m.date,
12     t.team_long_name AS awayteam, m.away_goal,
13   FROM match AS m
14   LEFT JOIN team AS t
15     ON m.awayTeam_id = t.team_api_id,
16     -- Select date, home_goal, and away_goal
17   SELECT
18     m.date,
19     hometeam,
20     awayteam,
21     m.home_goal,
22     m.away_goal,
23     -- Join away and home on the id column
24   FROM away
25   CROSS JOIN home
26   ON home.id = away.id;

```

query result

date	hometeam	awayteam	home_goal	away_goal
2011-07-29	Oud-Heverlee Leuven	RSC Anderlecht	2	1
2011-07-30	RAEC Mons	Standard de Liège	1	1
2011-07-30	KRC Genk	Beerschot AC	3	1
2011-07-30	KAAS Gent	KSV Cercle Brugge	0	1

Showing 100 out of 1000 rows

Windows Functions

What you've learned so far

- CASE statements
- Simple subqueries
- Nested and correlated subqueries
- Common table expressions

① ROW-NUMBERS() OVER()

② LAG(column, n) OVER(..) returns column's value at row n before current row

Current and last champions

Query

```
WITH Discus_Gold AS (
  SELECT
    Year, Country AS Champion
  FROM Summer_Medals
  WHERE
    Year IN (1996, 2000, 2004, 2008, 2012)
    AND Gender = 'Men' AND Medal = 'Gold'
    AND Event = 'Discus Throw')

SELECT
  Year, Champion,
  LAG(Champion, 1) OVER
  (ORDER BY Year ASC) AS Last_Champion
FROM Discus_Gold
ORDER BY Year ASC;
```

Result

Year	Champion	Last_Champion
1996	GER	null
2000	LTU	GER
2004	LTU	LTU
2008	EST	LTU
2012	GER	EST

③ PARTITION BY (similar like group-by
 but the results aren't
 rolled into one column)
 (ROW_NUMBER And
 LAG will Reset by partition)

Partitioning by one column

Query	Result																												
<pre>WITH Discus_Gold AS (...) SELECT Year, Event, Champion, LAG(Champion) OVER (PARTITION BY Event ORDER BY Event ASC, Year ASC) AS Last_Champ FROM Discus_Gold ORDER BY Event ASC, Year ASC;</pre>	<table border="1"> <thead> <tr> <th>Year</th> <th>Event</th> <th>Champion</th> <th>Last_Champion</th> </tr> </thead> <tbody> <tr> <td>2004</td> <td>Discus Throw</td> <td>LTU</td> <td>null</td> </tr> <tr> <td>2008</td> <td>Discus Throw</td> <td>EST</td> <td>LTU</td> </tr> <tr> <td>2012</td> <td>Discus Throw</td> <td>GER</td> <td>EST</td> </tr> <tr> <td>2004</td> <td>Triple Jump</td> <td>SWE</td> <td>null</td> </tr> <tr> <td>2008</td> <td>Triple Jump</td> <td>POR</td> <td>SWE</td> </tr> <tr> <td>2012</td> <td>Triple Jump</td> <td>USA</td> <td>POR</td> </tr> </tbody> </table>	Year	Event	Champion	Last_Champion	2004	Discus Throw	LTU	null	2008	Discus Throw	EST	LTU	2012	Discus Throw	GER	EST	2004	Triple Jump	SWE	null	2008	Triple Jump	POR	SWE	2012	Triple Jump	USA	POR
Year	Event	Champion	Last_Champion																										
2004	Discus Throw	LTU	null																										
2008	Discus Throw	EST	LTU																										
2012	Discus Throw	GER	EST																										
2004	Triple Jump	SWE	null																										
2008	Triple Jump	POR	SWE																										
2012	Triple Jump	USA	POR																										

POSTGRESQL SUMMARY STATS AND WINDOW FUNCTIONS

The four functions

Relative

- `LAG(column, n)` returns `column`'s value at the row `n` rows before the current row
- `LEAD(column, n)` returns `column`'s value at the row `n` rows after the current row

Absolute

- `FIRST_VALUE(column)` returns the first value in the table or partition
- `LAST_VALUE(column)` returns the last value in the table or partition

Partitioning with LEAD

- `LEAD(Champion, 1)` without `PARTITION BY`

Year	Event	Champion	Next_Champion
2004	Discus Throw	LTU	EST
2008	Discus Throw	EST	GER
2012	Discus Throw	GER	SWE
2004	Triple Jump	SWE	POR
2008	Triple Jump	POR	USA
2012	Triple Jump	USA	null

- `LEAD(Champion, 1)` with `PARTITION BY Event`

Year	Event	Champion	Next_Champion
2004	Discus Throw	LTU	EST
2008	Discus Throw	EST	GER
2012	Discus Throw	GER	null
2004	Triple Jump	SWE	POR
2008	Triple Jump	POR	USA
2012	Triple Jump	USA	null

So far, you've seen the functions used over the entire table, but what happens if you want to partition?

The ranking functions

- `ROW_NUMBER()` always assigns unique numbers, even if two rows' values are the same
- `RANK()` assigns the same number to rows with identical values, skipping over the next numbers in such cases
- `DENSE_RANK()` also assigns the same number to rows with identical values, but doesn't skip over the next numbers

Query

```
WITH Country_Medals AS (
  SELECT
    Country, COUNT(*) AS Medals
  FROM Summer_Medals
  GROUP BY Country),

  SELECT
    Country, Medals,
    NTILE(3) OVER (ORDER BY Medals DESC) AS Third
  FROM Country_Medals;
```

Result

Country	Medals	Third
USA	4585	1
URS	2849	1
GBR	1720	1
...
CZE	56	2
LTU	55	2
...
DOM	6	3
BWI	5	3
...

MAX Window function

Query

```
WITH Brazil_Medals AS (...)

SELECT
    Year, Medals,
    MAX(Medals)
    OVER (ORDER BY Year ASC) AS Max_Medals
FROM Brazil_Medals;
```

Result

Year	Medals	Max_Medals
1992	13	13
1996	5	13
2004	18	18
2008	14	18
2012	14	18

注意这里没有用range from function, 所以是根据已知的找最大

So far

SUM()

The screenshot shows a PostgreSQL course interface with two queries and their results.

Query:

```
WITH Medals AS (...)  
SELECT Year, Country, Medals,  
SUM(Medals) OVER (...)  
FROM Medals;
```

Result:

Year	Country	Medals	Medals_RT
2004	BRA	18	18
2008	BRA	32	32
2012	BRA	46	46
2004	CUB	31	31

Query:

```
WITH Medals AS (...)  
SELECT Year, Country, Medals,  
SUM(Medals) OVER (PARTITION BY Country ...)  
FROM Medals;
```

Result:

Year	Country	Medals	Medals_RT
2004	BRA	18	18
2008	BRA	14	32
2012	BRA	14	46
2004	CUB	31	31

A note in the result area states: "The source table is expanded to include a country column and Cuba's earned medals."

Min()

Frame :

- Frame: RANGE BETWEEN UNBOUNDED PRECEDING AND UNBOUNDED FOLLOWING
- Without the frame, LAST_VALUE would return the row's value in the City column
- By default, a frame starts at the beginning of a table or partition and ends at the current

ROWS BETWEEN

- ROWS BETWEEN [START] AND [FINISH]
 - n PRECEDING : n rows before the current row
 - CURRENT ROW : the current row
 - n FOLLOWING : n rows after the current row

Examples

- ROWS BETWEEN 3 PRECEDING AND CURRENT ROW
- ROWS BETWEEN 1 PRECEDING AND 1 FOLLOWING
- ROWS BETWEEN 5 PRECEDING AND 1 PRECEDING
 - 4 rows
 - 3 rows
 - 5 rows

Overview

- Moving average (MA): Average of last n periods
 - Example: 10-day MA of units sold in sales is the average of the last 10 days' sold units
 - Used to indicate momentum/trends
 - Also useful in eliminating seasonality
- Moving total: Sum of last n periods
 - Example: Sum of the last 3 Olympic games' medals
 - Used to indicate performance; if the sum is going down, overall performance is going down

ROWS vs RANGE

- RANGE BETWEEN [START] AND [FINISH]
 - Functions much the same as ROWS BETWEEN
 - RANGE treats duplicates in OVER's ORDER BY subclause as a single entity

Table

Year	Medals	Rows_RT	Range_RT
1992	10	10	10
1996	50	60	110
2000	50	110	110
2004	60	170	230
2008	60	230	230
2012	70	300	300

- ROWS BETWEEN is almost always used over RANGE BETWEEN

Pivoting
CROSSTAB

Enter CROSSTAB

```
CREATE EXTENSION IF NOT EXISTS tablefunc;
```

```
SELECT * FROM CROSSTAB($$  
source_sql TEXT  
$$) AS ct (column_1 DATA_TYPE_1,  
           column_2 DATA_TYPE_2,  
           ...,  
           column_n DATA_TYPE_N);
```

原sql语句

双引号，不能单引号

The screenshot shows a MySQL query editor with a complex query in the 'query' tab. The query uses multiple subqueries and joins to calculate gold medals won by countries in 2008. The results are displayed in the 'result' tab, showing three rows: CHN, RUS, and GBR.

Annotations in Chinese:

- 里面没有Country (Inside, there is no Country)
- 无双引号 (No double quotes)
- 有双引号 (Has double quotes)

Country	Medals	Year	Rank
CHN	3	2008	1
RUS	2	2008	2
GBR	1	2008	3

ROLLUP and CUBE

The old way

```

SELECT
    Country, Medal, COUNT(*) AS Awards
FROM Summer_Medals
WHERE
    Year = 2008 AND Country IN ('CHN', 'RUS')
GROUP BY Country, Medal
ORDER BY Country ASC, Medal ASC
UNION ALL

SELECT
    Country, 'Total', COUNT(*) AS Awards
FROM Summer_Medals
WHERE
    Year = 2008 AND Country IN ('CHN', 'RUS')
GROUP BY Country, 2
ORDER BY Country ASC;

```

↓ 更简洁.

Enter ROLLUP

```

SELECT
    Country, Medal, COUNT(*) AS Awards
FROM Summer_Medals
WHERE
    Year = 2008 AND Country IN ('CHN', 'RUS')
GROUP BY Country, ROLLUP(Medal)
ORDER BY Country ASC, Medal ASC;

```

- ROLLUP is a GROUP BY subclause that includes extra rows for group-level aggregations
- GROUP BY Country, ROLLUP(Medal) will count all Country - and Medal -level totals, then count only Country -level totals and fill in Medal with null's for these rows

ROLLUP - Query

```
SELECT
    Country, Medal, COUNT(*) AS Awards
FROM summer_medals
WHERE
    Year = 2008 AND Country IN ('CHN', 'RUS')
GROUP BY ROLLUP(Country, Medal)
ORDER BY Country ASC, Medal ASC;
```

- ROLLUP is hierarchical, de-aggregating from the leftmost provided column to the right-most
 - ROLLUP(Country, Medal) includes Country -level totals
 - ROLLUP(Medal, Country) includes Medal -level totals

Enter CUBE

```
SELECT
    Country, Medal, COUNT(*) AS Awards
FROM summer_medals
WHERE
    Year = 2008 AND Country IN ('CHN', 'RUS')
GROUP BY CUBE(Country, Medal)
ORDER BY Country ASC, Medal ASC;
```

- CUBE is a non-hierarchical ROLLUP
- It generates all possible group-level aggregations
 - CUBE(Country, Medal) counts Country -level, Medal -level, and grand totals

CUBE - Result

Country	Medal	Awards
CHN	Bronze	57
CHN	Gold	74
CHN	Silver	53
CHN	null	184
RUS	Bronze	56
RUS	Gold	43
RUS	Silver	44
RUS	null	143
null	Bronze	113
null	Gold	117
null	Silver	97
null	null	327

- Notice that Medal -level totals are included

ROLLUP vs CUBE

Enter COALESCE

- COALESCE() takes a list of values and returns the first non-`null` value, going from left to right
- `COALESCE(null, null, 1, null, 2) ? 1`
- Useful when using SQL operations that return `null`s
 - ROLLUP and CUBE
 - Pivoting
 - LAG and LEAD

Country	Medal	Awards
Both countries	All medals	327
CHN	All medals	184
CHN	Bronze	57
CHN	Gold	74
CHN	Silver	53
RUS	All medals	143
RUS	Bronze	56
RUS	Gold	43
RUS	Silver	44

Enter STRING_AGG

- `STRING_AGG(column, separator)` takes all the values of a column and concatenates them, with `separator` in between each value

`STRING_AGG(Letter, ', ',)` transforms this...

Letter
A
B
C

...into this

A, B, C

Sakila database